

How we can do things go faster. Concurrent and asynchronous technologies overview

Anton Mishchuk
Itransition, Minsk, 2012

The ultimate Ruby application

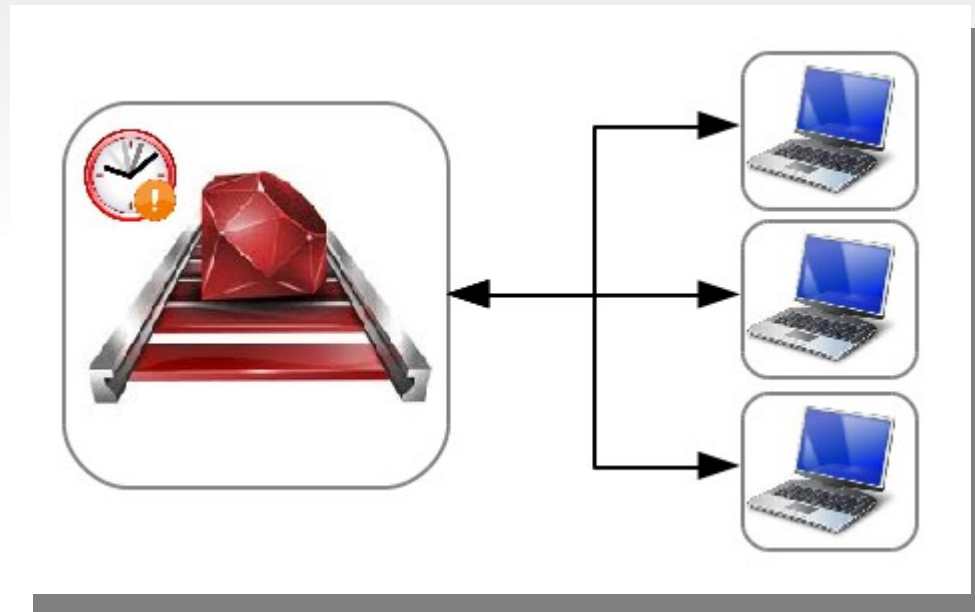


Contents

- Asynchronous execution of long tasks
- Reactor pattern and EventMachine
- HTML5 WebSocket
- Pub/sub systems and services

Few connections

Many calculations



Asynchronous executing long tasks in the background

Problems

- Disk manipulating
- Data processing
- Sending massive newsletters
- Statistics gathering
- Http downloads/uploads
- Warming caches
- Updating search index
- Slow insert statements



Asynchronous executing long tasks in the background Solutions

- Delayed_job (tobi, collectiveidea)
- Resque (defunkt)
- Beanstalkd with Stalker (kr, adamwiggins)



Main idea

- Store the tasks and execute them later in different process

Delayed_job

shopify.com - hosted ecommerce solution

- Queue storage – database (serialize Ruby objects)
- Supports many backends (ActiveRecord, DataMapper, Mongoid, Redis)
- Has job priorities
- Retries failed tasks
- Simple integration with Rails app



Delayed_job

- The best way:

```
class Work
  def perform
    #long running method
  end
end
Delayed::Job.enqueue Work.new("do it")
```

- Be carefull:

```
gem 'delayed_job_active_record'
class Calculation
  def calculate
    #long running method
  end
  handle_asynchronously :calculate
end
Calculation.new.calculate
```



Resque

github.com - social coding

- Inspired by DJ but designed for huge queues
- Based on Redis storage. Jobs are persisted as JSON objects
- Supports multiple queues
- Includes a Sinatra app for monitoring



```
class Worker
  @queue = :new_work
  def self.perform(work_id, params)
    work = Work.find(work_id)
    work.do_work(params)
  end
end

class Work
  def async_work(params)
    Resque.enqueue(Worker, self.id, params)
  end
end

Work.new.async_work('do_it')
console> QUEUE=:new_work rake resque:work
```

Resque monitor

Overview

Working

Failed

Queues

Workers

Stats


Queues

The list below contains all the registered queues with the number of jobs currently in the queue. Select a queue from above to view

Name	Jobs
<u>main</u>	6388
<u>failed</u>	0

1 of 1 Workers Working

The list below contains all workers which are currently running a job.

	Where	Queue	Processing
	<u>antonmi-1005PXD:6629</u>	MAIN	ResqueTask NaN days ago

Beanstalkd (with Stalker)

postrank.com - aggregator of social engagement

- Uses own memcached like storage
- Has priorities and timeouts
- Before and error tasks
- Simple & fast

```
class AnyClass
  def self.work(data)
    #long work with data
  end
end

Stalker.enqueue('stalker_job', data: id)

job "stalker_job" do |args|
  AnyClass.work(args[:data])
end
```



What is the best solution?

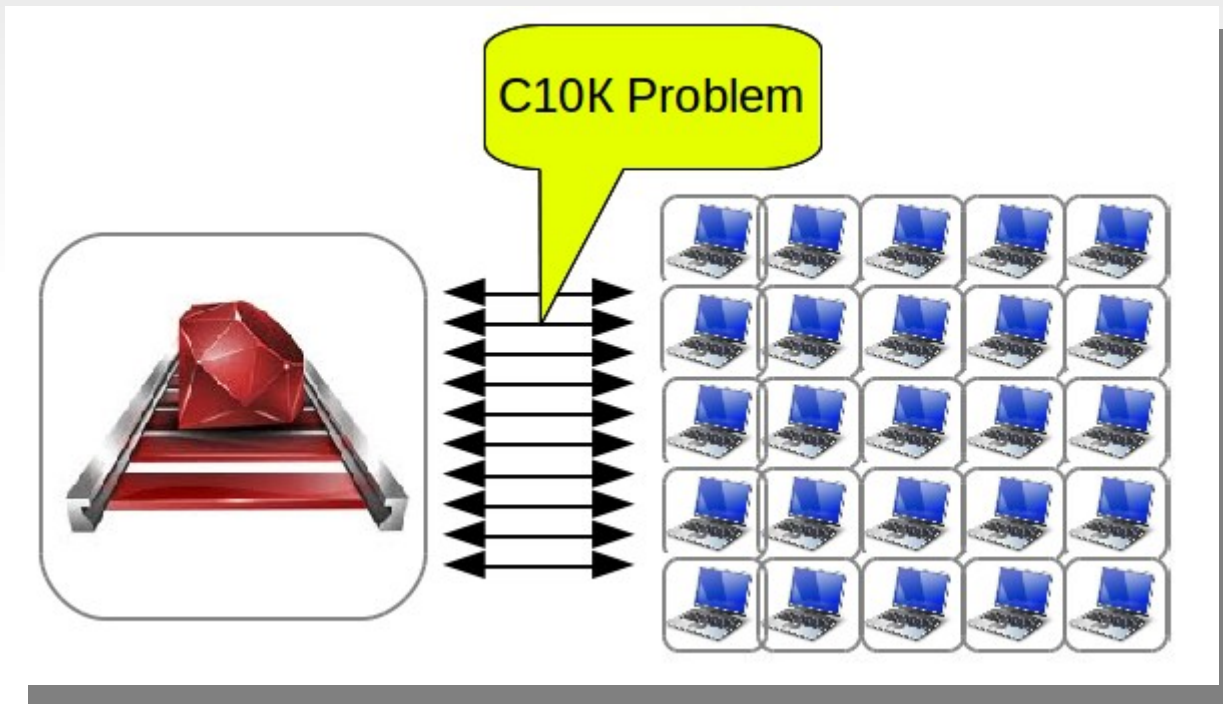
- It depends ...

	Enqueue, jobs/sec	Work, jobs/sec
delayed_job	200	120
resque	3800	300
beanstalkd	9000	5200

http://adam.heroku.com/past/2010/4/24/beanstalk_a_simple_and_fast_queueing_backend/

	Enqueue, 10000 tasks	Work, 10000 tasks
delayed_job	620 sec	1370 sec
resque	33 sec	310 sec
beanstalkd	6 sec	173 sec

Many connections Few calculations



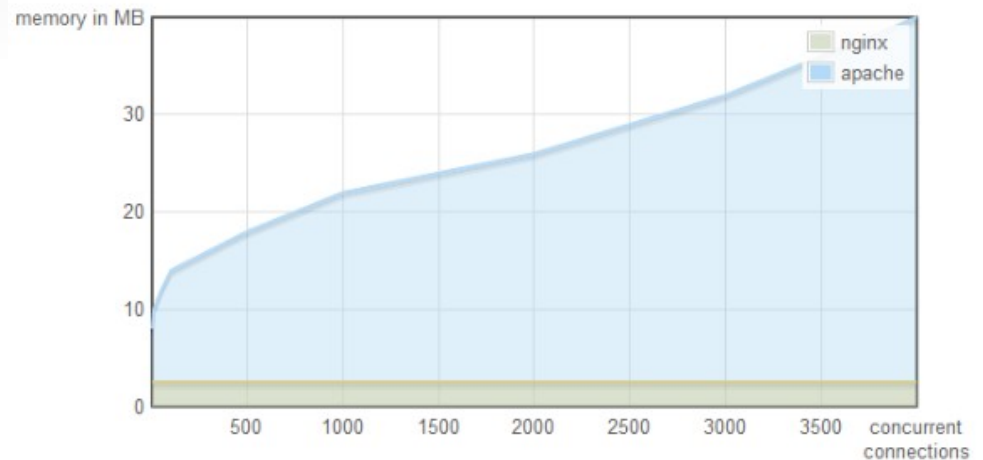
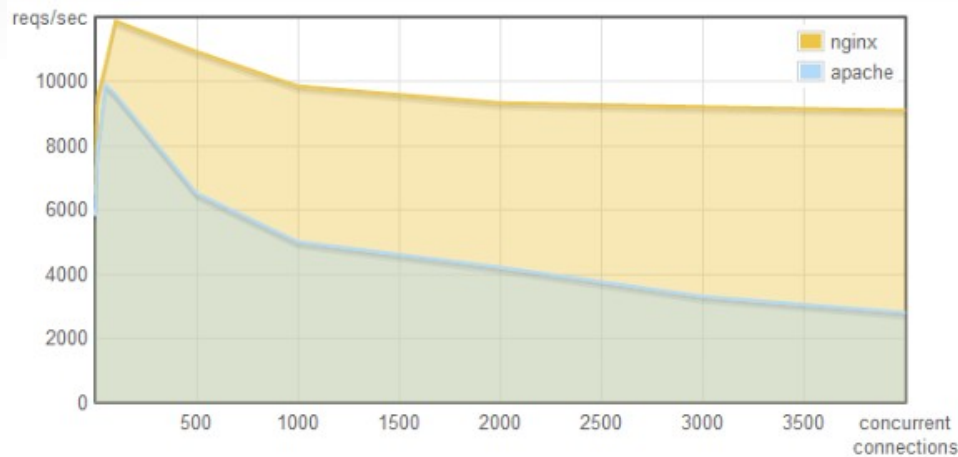
Threads or Events? Servers



Event Driven Based Server



Thread Based Server

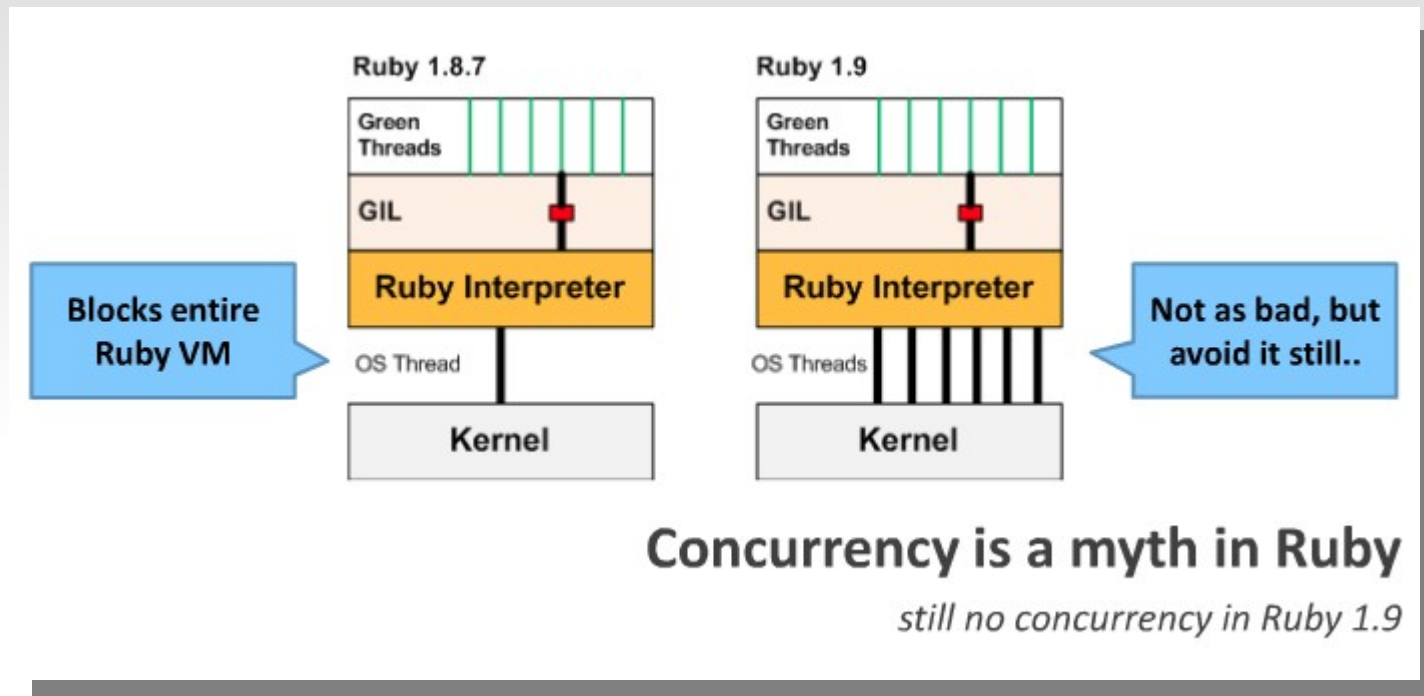


From Carlos Vasquez presentation:

http://www.slideshare.net/carlosforero3/ruby-eventmachine-emwebsocket?src=related_normal&rel=4675626



Threads in Ruby

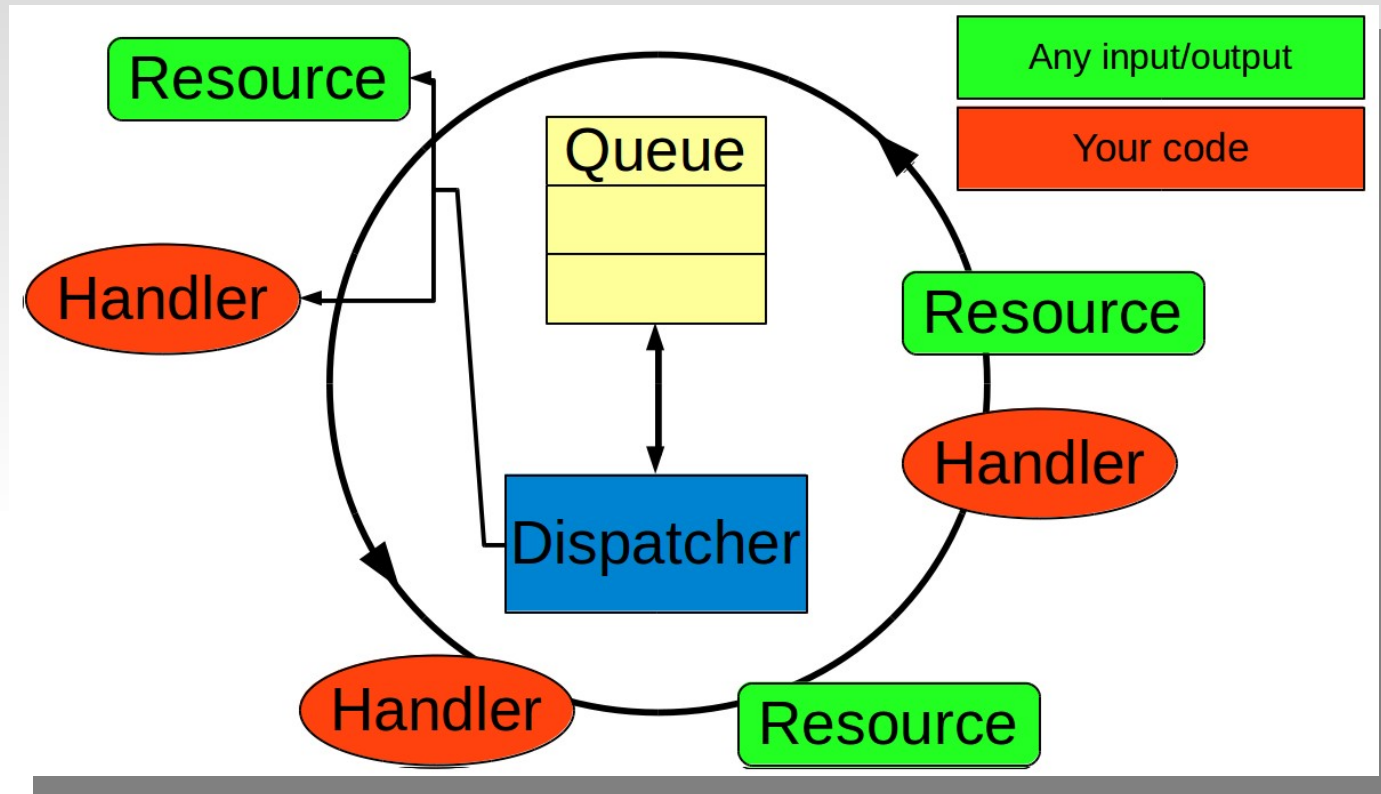


GIL (Global Interpreter Lock) is a mutual exclusion lock held by a programming language interpreter thread to avoid sharing code that is not thread-safe with other threads.

From Ilya Grigorik presentation:
<http://www.slideshare.net/igrigorik/no-callbacks-no-threads-railsconf-2010>



Reactor



- Reactor is a single threaded while loop ("reactor loop")
- The code in the loop "reacts" to incoming events
- If event handler takes too long, other events cannot fire
- Ideal for I/O based application

Reactor pattern implementations

- Ruby: EventMachine
- Javascript: Node.js
- Python: Twisted
- Java: Jboss Netty
- C#: Interlace
- ...

EventMachine

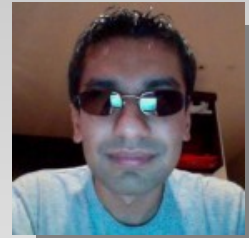


Interlace



EventMachine

tmm1 (Aman Gupta)



- Most web application are I/O bound, not CPU bound
- Basic idea – instead of waiting a response from the network use that time to process other request
- EventMachine provides event-driven I/O using the Reactor pattern

github

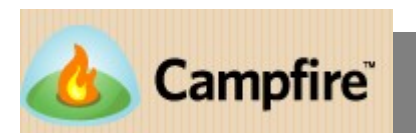
heroku

Thin



PostRank™

Goliath



EventMachine API

- EM.run
- EM.stop
- EM.next_tick
- EM.schedule
- EM.defer
- EM.system
- EM.connect
- EM.start_server
- EM.popen
- EM.watch
- EM::TickLoop
- EM::Deferrable
- EM.Callback
- EM::Timer
- EM::PeriodicTimer
- EM::Queue
- EM::Channel
- EM::Iterator
- EM::Connection
- EM::Protocols

Let's code

Protocol Implementation

- SMTP
 - HTTPClient
 - HTTPClient2
 - Postgres
 - MemCache
 - Stomp
 - Socks4
 - ObjectProtocol
 - SASLauth
 - LineAndText
 - LineText2
 -
- MySQL
 - Redis
 - Beanstalk
 - HTTPRequest
 - PubSubHubbub
 - Proxy
 - WebSocket
 - IRC
 - Cassandra
 - Solr
 - SSH
 -

Ilya Grigorik (igrigorik)



- Social & Google+ Analytics at Google
- Founder & CTO of PostRank (Google)
- Open-source evangelist
- Web engineer
- Photographer
- www.igvita.com
- em-http-request
- em-websocket
- em-synchrony
- etc, etc ...

em-http-request

- Keep-Alive and HTTP pipelining support
- Auto-follow 3xx redirects with max depth
- Automatic gzip & deflate decoding
- Streaming response processing
- Streaming file uploads
- HTTP proxy and SOCKS5 support
- Basic Auth & OAuth
- Connection-level & Global middleware support

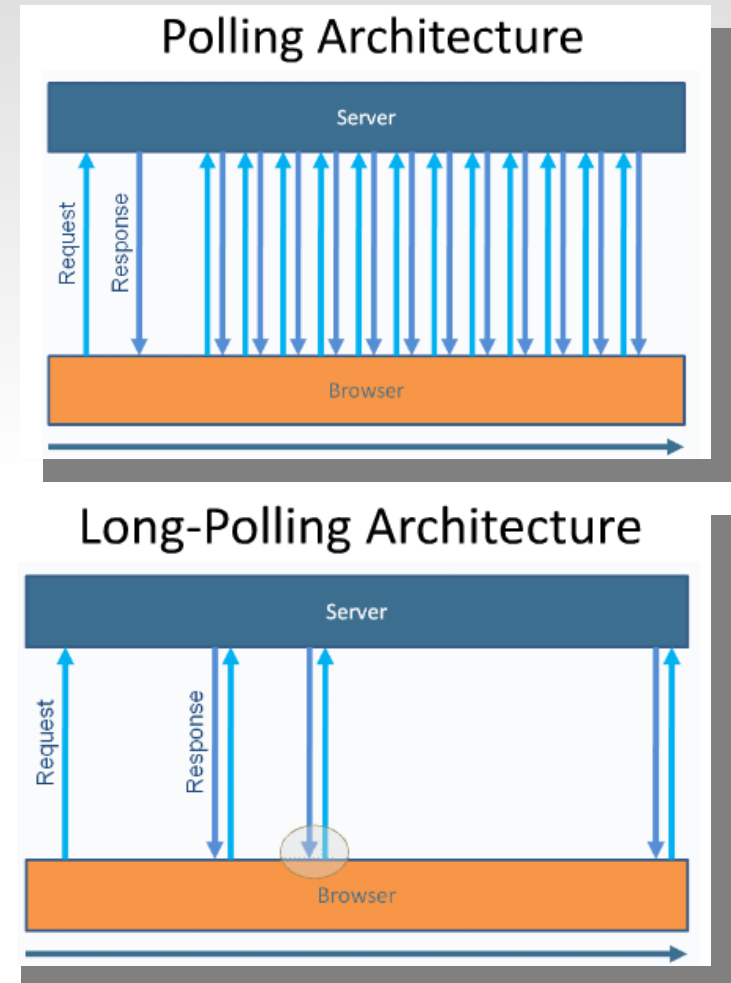
Real-time

- Notifications
- Chat
- Stocks
- Games
- Collaboration
- etc ...
- etc ...



Real-time?

- Polling (AJAX)
- Long Polling (Comet)



From ffdead presentaion: <http://www.slideshare.net/ffdead/the-html5-websocket-api>

HTTP overhead

```
GET /PollingStock//PollingStock HTTP/1.1
Host: localhost:8080
User-Agent: Mozilla/5.0 (Windows; U; Windows NT 5.1; en-US; rv:1.9.1.5)
Gecko/20091102 Firefox/3.5.5
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
Accept-Language: en-us
Accept-Encoding: gzip,deflate Accept-Charset: ISO-8859-1,utf-8;q=0.7,*;q=0.7
Keep-Alive: 300      691 chars
Connection: keep-alive
Referer: http://localhost:8080/PollingStock/
Cookie: showInheritedConstant=false;
       showInheritedProtectedConstant=false; showInheritedProperty=false;
       showInheritedProtectedProperty=false; showInheritedMethod=false;
       showInheritedProtectedMethod=false; showInheritedEvent=false; showInheritedStyle=false;
       showInheritedEffect=false
HTTP/1.x 200 OK
X-Powered-By: Servlet/2.5 Server: Sun Java System Application Server 9.1_02
Content-Type: text/html; charset=UTF-8
Content-Length: 321
Date: Sat, 07 Nov 2009 00:32:46 GMT
```

Total 871 chars

1.000 clients polling every second:

$$1.000 \cdot 871 \cdot 8 = 7 \text{ Mbps !!!!!}$$

Websocket



- WebSocket protocol - connection established by upgrading from HTTP to WebSocket protocol
- True full-duplex communication channel
- Proxy/Firewall friendly:
 - runs via port 80/443
 - integrates with cookie based authentication
- Secure connection via Secure WebSockets

Websocket

- Each message has only 2 bytes of overhead (0x00<data>0xFF)
- No latency from establishing new connection
- No polling overhead (only sends messages when there is something to send)

1.000 clients send message every second:

$$1.000 \cdot 2 \cdot 8 = 16 \text{ kbps (was } 7 \text{ Mbps)}$$



WebSocket API

- onopen
- onclose
- onmessage
- onerror
- send
- close

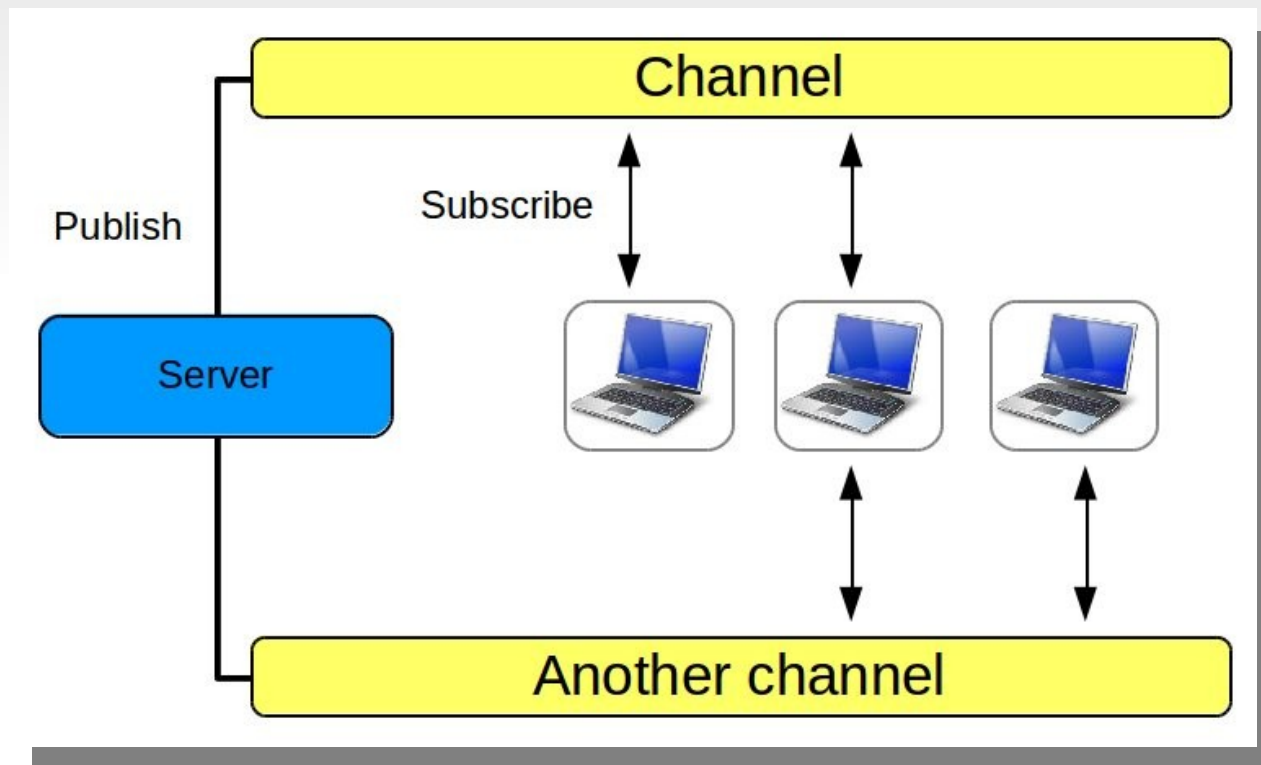
em-websocket

- See example

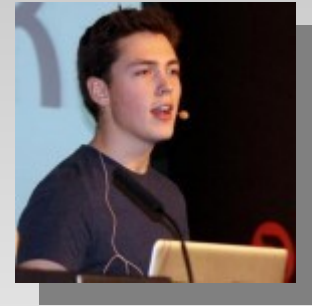
EventMachine in Rails

- Run EM in a thread
- Use EM enabled webserver:
Thin, Goliath, Rainbow
- See example

Publish Subscribe pattern



Juggernaut (maccman)



- Node.js server
- Redis storage
- Ruby client
- Supports the following protocols:
WebSocket, Adobe Flash Socket, ActiveX HTMLFile , XHR, ...
- Horizontal scaling
- Reconnection support
- SSL support

Faye (jcoglan)



- Node.js or EventMachine servers
 - Bayeux protocol (primarily over HTTP)
 - Memory or Redis storage
 - Faye::WebSocket
 - etc, etc ...
-
- private_pub gem (Ryan Bates)



Private Pub

- In the view:

```
PrivatePub.subscribe("/messages", function(data, channel){  
  //do smth with data  
})
```

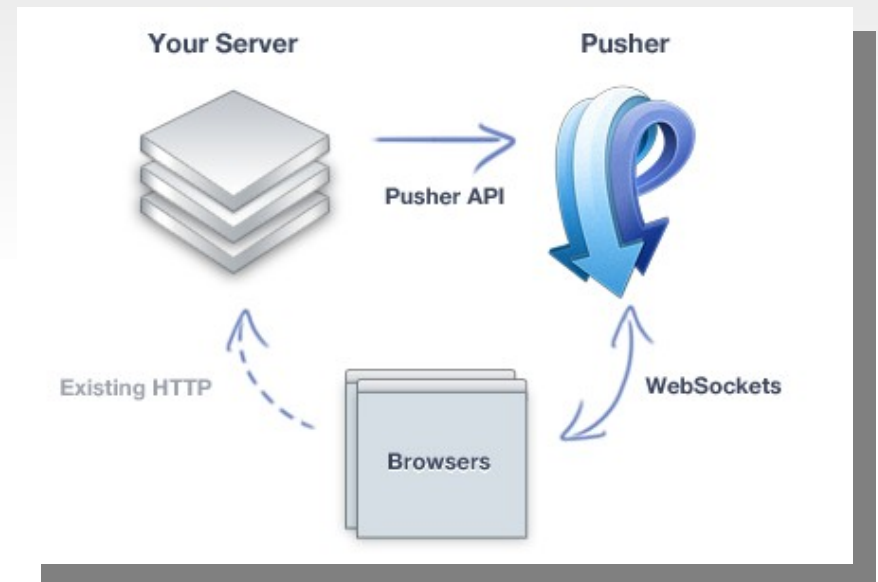
- In the controller:

```
PrivatePub.publish_to "/messages/", @message
```

Pusher.com

realtime messaging platform

- Pub/Sub model with channels, events and webhooks
- Rich suite of libraries
- WebSockets with fallback to Flash
- Public, private and presence channel
- Tools for monitoring and debugging



Pusher client libraries

- Javascript
- iOS – Objective C
- iOS – Appcelerator Titanium
- Android – Java
- ActionScript
- .Net & Silverlight
- Ruby
- Ardunio

Pusher publisher libraries

- Node.js
- Java
- Groovy/Grails
- Clojure
- Python
- Ruby
- VB.NET
- C#
- PHP
- Perl
- Coldfushion

Pusher Channels

- Public channels – can be subscribed by anyone.
- Private channels – allow controll access to the data you are broadcasting
- Presence channels – let you register user information on subscription and let other members of channel know who is online

Pusher events

- Connection events: connecting, connected, disconnected, unavailable, failed
- Channel events: subscription_succeeded, subscription_error, member_added, member_removed
- Custom events for channel(s)

```
var pusher = new Pusher('API_KEY');  
var channel = pusher.subscribe('APPL');  
channel.bind('new-price',  
  function(data) {  
    // add new price into the APPL widget  
  }  
);
```

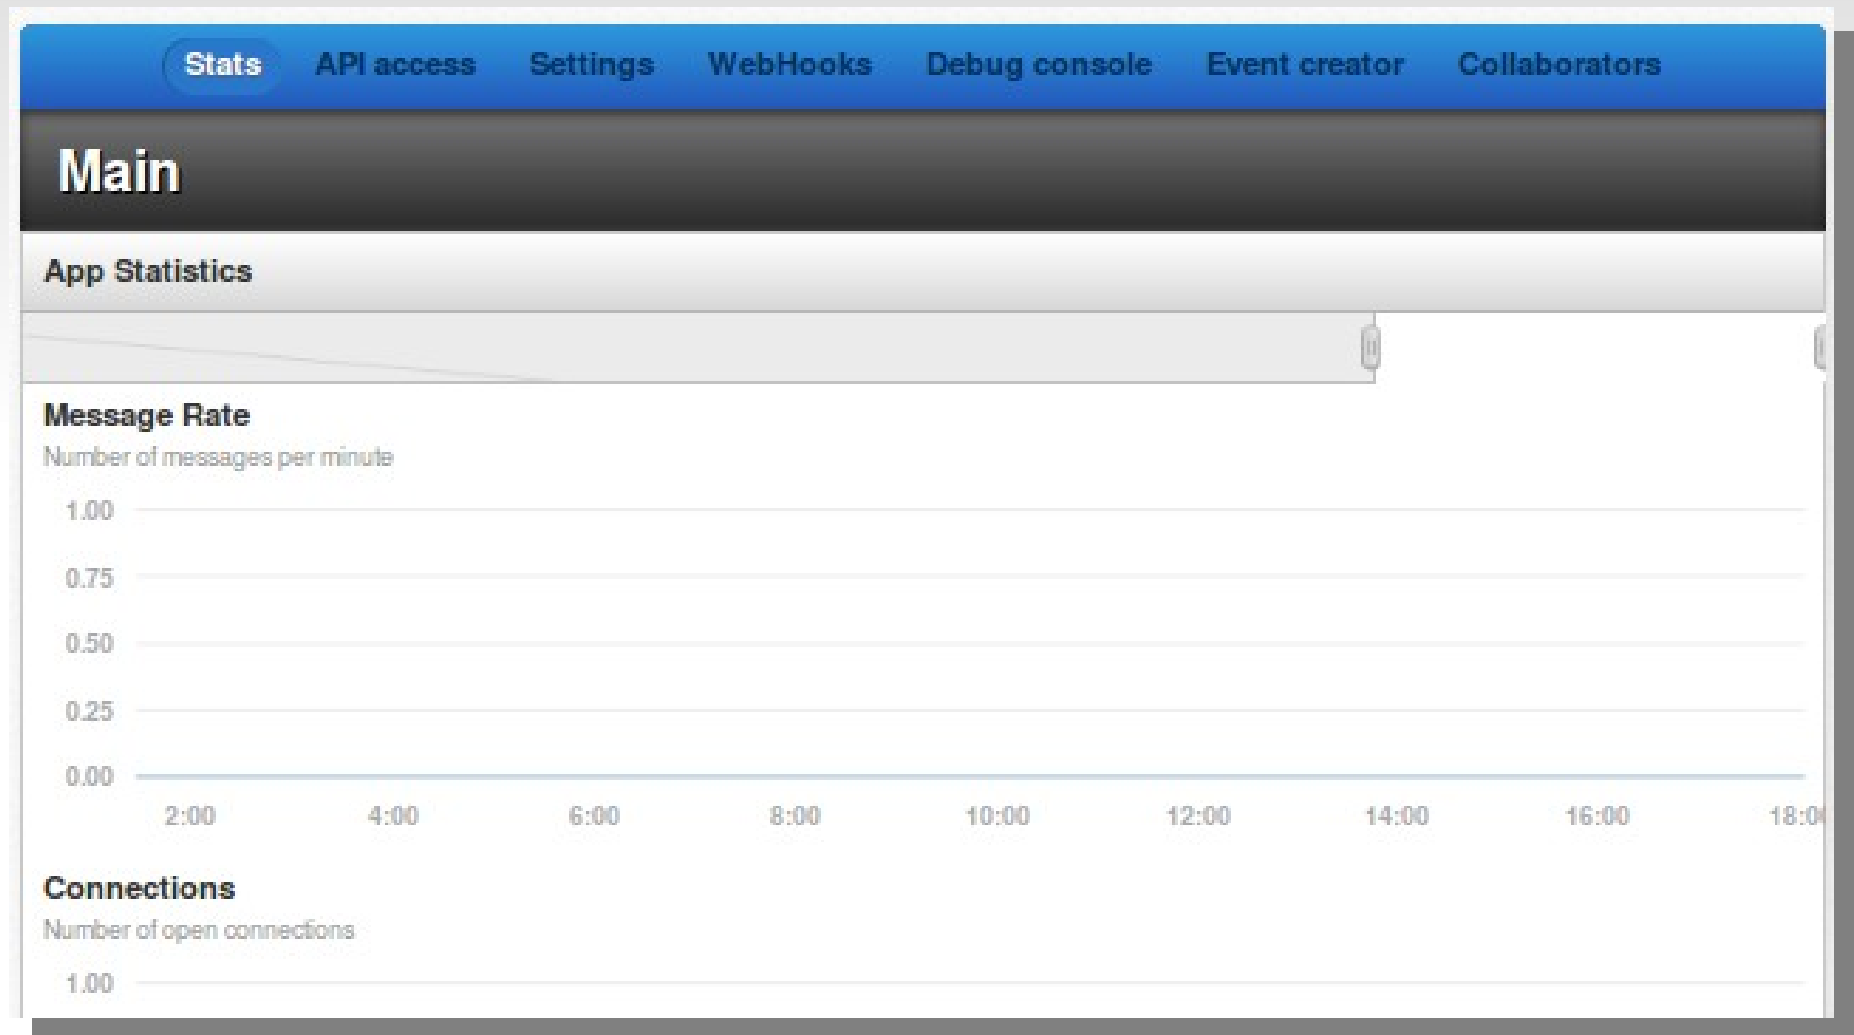
Pusher webhooks

- Allow your server to be notified about events occurring within Pusher
- Pusher sends HTTP POST request to the specified url

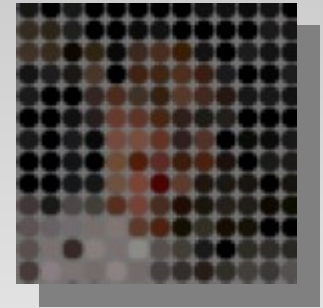
```
{  
  "time_ms": 1327078148132  
  "events": [  
    { "name": "event_name", "some": "data" }  
  ]  
}
```

- Examples: channel_occupied, channel_vacated

Pusher monitoring and debugging



pusher-gem (mlaughran)



- Configure app

```
Pusher.app_id = '16661'  
Pusher.key = '7382735ce49c18967856'  
Pusher.secret = '1d767e0dc26bcbd11329'
```

```
<script type="text/javascript">  
  var pusher = new Pusher('7382735ce49c18967856');  
  var demo_channel = pusher.subscribe('demo_channel');  
  demo_channel.bind('show_message', function(data){  
    console.log(data);  
    $('#message').html(data);  
  })  
</script>
```

- Trigger events

```
Pusher['demo_channel'].trigger('show_message', "Hello Pusher")
```

- Use EM

```
def pusher  
  EM.schedule do  
    EM::Timer.new(5) do  
      push = Pusher['demo_channel'].trigger_async('show_message', "Use EM")  
      push.callback {puts "Message sent"}  
    end  
  end  
end
```

Pusher Pricing

Sandbox	Bootstrap	Startup	Big Boy	Enterprise
FREE	\$19/month	\$49/month	\$199/month	Contact us
20 Max Connections	100 Max Connections	500 Max Connections	5,000 Max Connections	Millions Connections
100,000 Messages per day	200,000 Messages per day	1 million Messages per day	10 million Messages per day	Billions Messages per day
Encryption SSL Protection	Encryption SSL Protection	Encryption SSL Protection	Encryption SSL Protection	Premium support 24/7 Telephone

Defend Ruby from "Ruby is slow" people



References

Delayed Job and others

- https://github.com/tobi/delayed_job
- https://github.com/collectiveidea/delayed_job
- <http://railscasts.com/episodes/171-delayed-job-revised>
- <https://github.com/defunkt/resque>
- <http://railscasts.com/episodes/271-resque>
- <http://kr.github.com/beanstalkd/>
<https://github.com/adamwiggins/stalker>
- <http://railscasts.com/episodes/243-beanstalkd-and-stalker>
- http://adam.heroku.com/past/2010/4/24/beanstalk_a_simple_and_fast_queue

References

EventMachine

- <https://github.com/eventmachine/eventmachine>
- <http://www.youtube.com/watch?v=mPDs-xQhPb0>
- <http://www.viddler.com/v/cfadc37f>
- <https://github.com/igrigorik/em-synchrony>
- <http://www.igvita.com/2010/03/22/untangling-evented-code-with-ruby-fibers/>
- <http://www.igvita.com/2009/12/22/ruby-websockets-tcp-for-the-browser/>
- <https://github.com/igrigorik/async-rails>
- <http://www.slideshare.net/jweiss/eventmachine>
- <http://www.slideshare.net/igrigorik/no-callbacks-no-threads-railsconf-2010>
- <http://www.slideshare.net/autonomous/ruby-concurrency-and-eventmachine>

References

EventMachine

- http://www.slideshare.net/carlosforero3/ruby-eventmachine-emwebsocket?src=related_normal&rel=481
- http://www.slideshare.net/kbal11/ruby-19-fibers?src=related_normal&rel=481
- <http://www.slideshare.net/KyleDrake/fast-concurrent-ruby-web-applications-w>
- <http://www.scribd.com/doc/28253878/EventMachine-scalable-non-blocking-i-c>
- <http://code.macournoyer.com/thin/>
- <http://www.igvita.com/2011/03/08/goliath-non-blocking-ruby-19-web-server/>
- <http://www.slideshare.net/ismasan/websockets-and-ruby-eventmachine>
- <http://www.slideshare.net/ffdead/the-html5-websocket-api>

References

Pub/Sub

- <http://juggernaut.rubyforge.org/>
- <https://github.com/maccman/juggernaut>
- <http://faye.jcoglan.com/>
- <http://railscasts.com/episodes/260-messaging-with-faye>
- https://github.com/ryanb/private_pub
- <http://railscasts.com/episodes/316-private-pub>
- <https://github.com/pusher/pusher-gem/tree/master/examples>
- <http://pusher.com>
- <https://github.com/pusher/pusher-gem>

Thank you

- Questions?